

# End-to-End PostgreSQL Data Architecture: Partitioning, Query Optimization, and ETL Automation at Scale

-- Keywords: data engineering, query optimization, window functions, CTE, partitioning, indexing,

-- materialized view, JSONB, full-text search, concurrency, CDC, incremental ETL, schema evolution, upsert

-- DB: PostgreSQL (12+), adjust dialect-specific parts for other systems

---

-- 0) Safety: Work in a sandbox schema

---

DROP SCHEMA IF EXISTS prod\_analytics CASCADE;

CREATE SCHEMA prod\_analytics AUTHORIZATION CURRENT\_USER;

SET search\_path = prod\_analytics, public;

---

-- 1) Core OLTP table: users & organizations (demonstrates FK, constraints, generated cols)

---

CREATE TABLE prod\_analytics.organizations (

    org\_id        BIGSERIAL PRIMARY KEY,

    org\_name      TEXT NOT NULL,

    industry      TEXT,

```
    created_at    TIMESTAMP WITH TIME ZONE DEFAULT now(),
    metadata      JSONB DEFAULT '{}':JSONB
);

COMMENT ON TABLE prod_analytics.organizations IS 'Organization master (data
engineering, lineage metadata)';
```

```
CREATE TABLE prod_analytics.users (
    user_id       BIGSERIAL PRIMARY KEY,
    org_id        BIGINT NOT NULL REFERENCES prod_analytics.organizations(org_id) ON
DELETE CASCADE,
    email         CITEXT NOT NULL UNIQUE,
    username      TEXT,
    signup_ts     TIMESTAMP WITH TIME ZONE DEFAULT now(),
    attrs         JSONB DEFAULT '{}':JSONB,
    -- generated column: domain extracted from email for analytics segmentation
    email_domain  TEXT GENERATED ALWAYS AS (split_part(email, '@', 2)) STORED,
    is_active     BOOLEAN DEFAULT TRUE,
    last_login    TIMESTAMP WITH TIME ZONE,
    created_at    TIMESTAMP WITH TIME ZONE DEFAULT now()
);

CREATE INDEX idx_users_org_created ON prod_analytics.users (org_id, created_at DESC);

COMMENT ON TABLE prod_analytics.users IS 'User master table with generated column for
email_domain';
```

-----

-- 2) Event table for analytics (append-only, partitioned by day, optimized for high ingest)

-- Uses RANGE partitioning for time-series scale; BRIN index for fast scans on large data.

-----

```
CREATE TABLE prod_analytics.events (  
    event_id      BIGSERIAL NOT NULL,  
    user_id       BIGINT REFERENCES prod_analytics.users(user_id),  
    org_id        BIGINT,  
    event_type     TEXT NOT NULL,  
    event_props    JSONB,  
    event_time     TIMESTAMP WITH TIME ZONE NOT NULL DEFAULT now(),  
    received_at    TIMESTAMP WITH TIME ZONE DEFAULT now(),  
    -- useful derived column for micro-bucketing (example)  
    event_date     DATE GENERATED ALWAYS AS (event_time::date) STORED  
) PARTITION BY RANGE (event_date);
```

-- create monthly partitions for 12 months (example)

DO \$\$

BEGIN

FOR i IN 0..11 LOOP

EXECUTE format('

CREATE TABLE IF NOT EXISTS prod\_analytics.events\_p%s PARTITION OF  
prod\_analytics.events

FOR VALUES FROM (date\_trunc("month", current\_date) + INTERVAL "%s month")

TO (date\_trunc("month", current\_date) + INTERVAL "%s month")

', i, i, i+1);

END LOOP;

```
END $$;
```

```
-- Indexing: composite btree for common filters; BRIN for very large time scans;
```

```
CREATE INDEX IF NOT EXISTS idx_events_user_event_time ON prod_analytics.events  
(user_id, event_time DESC);
```

```
CREATE INDEX IF NOT EXISTS idx_events_org_time ON prod_analytics.events (org_id,  
event_time DESC);
```

```
-- BRIN index on partitioned table (benefit for huge append-only data)
```

```
CREATE INDEX IF NOT EXISTS brin_events_event_date ON prod_analytics.events USING  
BRIN (event_date);
```

```
COMMENT ON TABLE prod_analytics.events IS 'Append-only event table partitioned by date.  
Keywords: partitioning, BRIN, time-series.';
```

```
-----  
-- 3) Full-text search & semantic tagging using tsvector and GIN (search across articles/support  
logs)  
-----
```

```
CREATE TABLE prod_analytics.knowledge_base (  
    kb_id          BIGSERIAL PRIMARY KEY,  
    org_id         BIGINT REFERENCES prod_analytics.organizations(org_id),  
    title          TEXT NOT NULL,  
    content        TEXT NOT NULL,  
    content_json   JSONB DEFAULT '{}':JSONB,  
    created_at     TIMESTAMP WITH TIME ZONE DEFAULT now(),  
    updated_at     TIMESTAMP WITH TIME ZONE DEFAULT now(),  
    -- materialized searchable vector (for full-text search)
```

```

search_vector TSVECTOR GENERATED ALWAYS AS (
    setweight(to_tsvector('english', coalesce(title,"")), 'A') ||
    setweight(to_tsvector('english', coalesce(content,"")), 'B')
) STORED
);

CREATE INDEX idx_kb_search_vector ON prod_analytics.knowledge_base USING GIN
(search_vector);

COMMENT ON TABLE prod_analytics.knowledge_base IS 'Knowledge base with full-text
search support (GIN).';

```

---

```

-- 4) Upsert pattern (ON CONFLICT) for idempotent ETL / CDC ingestion

```

---

```

-- Example staging table for CDC

```

```

CREATE TABLE prod_analytics.staging_users (
    src_id      TEXT PRIMARY KEY,
    email       CITEXT,
    username    TEXT,
    org_src     TEXT,
    payload     JSONB,
    ingested_at TIMESTAMP WITH TIME ZONE DEFAULT now()
);

```

```

-- Upsert into users from staging: idempotent CDC pattern

```

```

-- demonstrates upsert strategy, conflict_target, and returning clause

```

```

WITH upsert AS (

```

```

INSERT INTO prod_analytics.users (org_id, email, username, attrs, created_at)

SELECT o.org_id, s.email, s.username, s.payload, now()

FROM prod_analytics.staging_users s

LEFT JOIN LATERAL (

    SELECT org_id FROM prod_analytics.organizations WHERE org_name = s.org_src
LIMIT 1

) o ON TRUE

ON CONFLICT (email) DO UPDATE

SET username = EXCLUDED.username,

    attrs = users.attrs || EXCLUDED.attrs,

    last_login = CASE WHEN EXCLUDED.attrs->>'last_login' IS NOT NULL THEN
(EXCLUDED.attrs->>'last_login')::timestampz ELSE users.last_login END

RETURNING users.user_id, users.email, users.username

)

SELECT 'UPSERT_COMPLETE' as result, count(*) FROM upsert;

```

-----

-- 5) Recursive CTE: org hierarchy / reporting chain analytics

-- Demonstrates recursive CTEs and graph traversal

-----

```

CREATE TABLE prod_analytics.org_hierarchy (

    node_id    BIGSERIAL PRIMARY KEY,

    org_id     BIGINT REFERENCES prod_analytics.organizations(org_id),

    parent_id  BIGINT REFERENCES prod_analytics.org_hierarchy(node_id),

    role       TEXT,

    metadata   JSONB

```

```
);
```

```
-- Recursive traversal to get full path and depth
```

```
WITH RECURSIVE hierarchy AS (
```

```
    SELECT node_id, org_id, parent_id, role, metadata, 1 AS depth, array[node_id] AS path
```

```
    FROM prod_analytics.org_hierarchy WHERE parent_id IS NULL
```

```
    UNION ALL
```

```
    SELECT h.node_id, h.org_id, h.parent_id, h.role, h.metadata, r.depth + 1, r.path || h.node_id
```

```
    FROM prod_analytics.org_hierarchy h
```

```
    JOIN hierarchy r ON h.parent_id = r.node_id
```

```
)
```

```
SELECT * FROM hierarchy ORDER BY depth DESC LIMIT 100;
```

```
-----  
-- 6) Analytical query: rolling metrics and retention using window functions
```

```
-- Common MAANG interview prompt: DAU/MAU, retention, event funnel
```

```
-----  
-- DAU (daily active users) and 7-day rolling active users:
```

```
WITH dau AS (
```

```
    SELECT event_date,
```

```
        count(DISTINCT user_id) AS dau_count
```

```
    FROM prod_analytics.events
```

```
    WHERE event_time >= now() - INTERVAL '90 days'
```

```
    GROUP BY event_date
```

```
)
```

```
SELECT
    event_date,
    dau_count,
    SUM(dau_count) OVER (ORDER BY event_date ROWS BETWEEN 6 PRECEDING AND
CURRENT ROW) AS rolling_7_day_active_users
FROM dau
ORDER BY event_date DESC
LIMIT 90;
```

-- 7-day retention cohort analysis: cohort by signup\_date, retention by event\_date

```
WITH cohorts AS (
    SELECT u.user_id, u.created_at::date AS signup_date, e.event_time::date AS event_date
    FROM prod_analytics.users u
    LEFT JOIN prod_analytics.events e ON e.user_id = u.user_id
    WHERE u.created_at >= current_date - INTERVAL '90 days'
)
```

```
SELECT
    signup_date,
    event_date,
    COUNT(DISTINCT user_id) AS active_users,
    (event_date - signup_date) AS days_after_signup
FROM cohorts
GROUP BY signup_date, event_date
ORDER BY signup_date DESC, days_after_signup
LIMIT 200;
```



-----  
-- 7) Lateral join and JSONB-heavy query: extract nested properties & flatten arrays  
-----

-- Example event\_props has schema {"product": {"id": "...", "attributes": [{"k": "", "v": ""}]}}

```
SELECT e.event_id, e.user_id, e.event_time,  
       (e.event_props->'product'->>'id')::text AS product_id,  
       k.attr->>'k' AS attr_key,  
       k.attr->>'v' AS attr_value  
FROM prod_analytics.events e  
  
CROSS JOIN LATERAL jsonb_array_elements(coalesce(e.event_props->'product'->'attributes',  
['']::jsonb)) AS k(attr)  
  
WHERE e.event_type = 'product_view'  
  
LIMIT 500;
```

-----  
-- 8) Materialized View for expensive aggregated reports + incremental refresh strategy  
-----

```
CREATE MATERIALIZED VIEW IF NOT EXISTS prod_analytics.mv_org_daily_metrics  
AS  
  
SELECT  
    org_id,  
    event_date,  
    count(DISTINCT user_id) AS dau,  
    count(*) FILTER (WHERE event_type = 'purchase') AS purchases,
```

```
sum((event_props->>'price')::numeric) FILTER (WHERE (event_props->>'price') IS NOT
NULL) AS revenue,
```

```
now() AS computed_at
```

```
FROM prod_analytics.events
```

```
GROUP BY org_id, event_date;
```

```
-- For incremental refresh: using CONCURRENTLY (note: requires unique index on mat view)
```

```
CREATE UNIQUE INDEX IF NOT EXISTS mv_org_daily_metrics_idx ON
prod_analytics.mv_org_daily_metrics (org_id, event_date);
```

```
-- Refresh concurrently in maintenance job (example):
```

```
-- REFRESH MATERIALIZED VIEW CONCURRENTLY
prod_analytics.mv_org_daily_metrics;
```

```
-----
```

```
-- 9) PL/pgSQL: idempotent incremental ETL function + observability
```

```
-- Demonstrates transactional control, exception handling, logging, and metrics table
```

```
-----
```

```
CREATE TABLE IF NOT EXISTS prod_analytics.etl_jobs (
```

```
    job_id    BIGSERIAL PRIMARY KEY,
```

```
    job_name  TEXT NOT NULL,
```

```
    started_at TIMESTAMPTZ,
```

```
    finished_at TIMESTAMPTZ,
```

```
    status    TEXT,
```

```
    rows_processed BIGINT,
```

```
    details    JSONB
```

```
);
```

```

-- simple incremental ETL: move from events staging into partitioned events table

CREATE OR REPLACE FUNCTION prod_analytics.etl_ingest_events(batch_size INT
DEFAULT 10000)

RETURNS JSONB

LANGUAGE plpgsql

AS $$

DECLARE

    v_start TIMESTAMPTZ := now();

    v_rows BIGINT := 0;

    v_job_id BIGINT;

BEGIN

    INSERT INTO prod_analytics.etl_jobs (job_name, started_at, status)

    VALUES ('ingest_events', v_start, 'RUNNING') RETURNING job_id INTO v_job_id;

    -- Simulated idempotent move: assume staging_events exists with unique event_id

    WITH moved AS (

        DELETE FROM prod_analytics.staging_events se

        WHERE se.event_time < now() + INTERVAL '1 second'

        RETURNING se.*

    )

    INSERT INTO prod_analytics.events (user_id, org_id, event_type, event_props, event_time,
received_at)

    SELECT user_id, org_id, event_type, payload, event_time, now()

    FROM moved

    ON CONFLICT (event_id) DO NOTHING

```

```

RETURNING 1 INTO v_rows;

-- update job record

UPDATE prod_analytics.etl_jobs

    SET finished_at = now(), status = 'SUCCESS', rows_processed = COALESCE(v_rows,0)

WHERE job_id = v_job_id;

RETURN jsonb_build_object('job_id', v_job_id, 'rows', COALESCE(v_rows,0),
'duration_seconds', EXTRACT(EPOCH FROM now() - v_start));

EXCEPTION WHEN OTHERS THEN

    UPDATE prod_analytics.etl_jobs

        SET finished_at = now(), status = 'FAILED', details = jsonb_build_object('error', SQLERRM)

    WHERE job_id = v_job_id;

    RAISE;

END;

$$;

```

---

```

-- 10) Concurrency & advisory locks: ensure single-run for ETL, lock-based coordination

```

---

```

-- Use advisory locks for cross-process coordination:

-- SELECT pg_try_advisory_lock(hashtext('etl_ingest_events'));

-- run function

-- SELECT prod_analytics.etl_ingest_events();

```

---

-- 11) Observability & introspection queries: query plan + statistics

-----

-- Show top 10 slowest queries (requires pg\_stat\_statements extension). This is an example query; extension must be enabled.

-- (In a real environment, enable extension and run periodically)

CREATE EXTENSION IF NOT EXISTS pg\_stat\_statements;

-- Top queries by total\_time

SELECT queryid, query, calls, total\_time, mean\_time, rows

FROM pg\_stat\_statements

ORDER BY total\_time DESC

LIMIT 10;

-- Explain analyze template: show how to capture query plan metadata (demonstration only)

-- EXPLAIN (ANALYZE, BUFFERS, VERBOSE) SELECT \* FROM prod\_analytics.events  
WHERE event\_time > now() - INTERVAL '1 day';

-----

-- 12) Advanced indexing: partial, expression, bloom, GIN for JSONB & full-text

-----

-- partial index for active users with recent activity

CREATE INDEX IF NOT EXISTS idx\_users\_active\_recent ON prod\_analytics.users  
(last\_login) WHERE is\_active = TRUE;

-- expression index for lower(username) lookups

CREATE INDEX IF NOT EXISTS idx\_users\_username\_lower ON prod\_analytics.users  
((lower(username)));

-- GIN for JSONB path existence/containment

```
CREATE INDEX IF NOT EXISTS idx_events_props_gin ON prod_analytics.events USING
GIN (event_props jsonb_path_ops);
```

-- trigram index for fuzzy match (pg\_trgm extension)

```
CREATE EXTENSION IF NOT EXISTS pg_trgm;
```

```
CREATE INDEX IF NOT EXISTS idx_kb_title_trgm ON prod_analytics.knowledge_base
USING gin (title gin_trgm_ops);
```

-- BRIN index already created for event\_date earlier - cheap for large time series tables.

-----  
-- 13) Data quality checks (DQ) & assertions: example SQL tests  
-----

-- 13.1 Referential integrity check: events.user\_id that have no user in users table

```
SELECT e.user_id, count(*) AS orphan_event_count
FROM prod_analytics.events e
LEFT JOIN prod_analytics.users u ON u.user_id = e.user_id
WHERE u.user_id IS NULL
GROUP BY e.user_id
ORDER BY orphan_event_count DESC
LIMIT 50;
```

-- 13.2 Schema validation: expected keys present in event\_props for 'purchase' events

```

SELECT

count(*) FILTER (WHERE (event_props ? 'order_id') IS NOT TRUE) AS missing_order_id,

count(*) FILTER (WHERE (event_props ? 'price') IS NOT TRUE) AS missing_price,

count(*) AS total_purchase_events

FROM prod_analytics.events

WHERE event_type = 'purchase';

```

-----

-- 14) Example complex Business Intelligence query: funnel conversion & attribution

-----

-- Attribution: count unique users who viewed product -> added\_to\_cart -> purchased within 7 days

```

WITH views AS (

    SELECT user_id, event_time, (event_props->>'product_id') AS product_id

    FROM prod_analytics.events WHERE event_type = 'product_view'

),

adds AS (

    SELECT user_id, event_time, (event_props->>'product_id') AS product_id

    FROM prod_analytics.events WHERE event_type = 'add_to_cart'

),

purchases AS (

    SELECT user_id, event_time, (event_props->>'product_id') AS product_id

    FROM prod_analytics.events WHERE event_type = 'purchase'

)

SELECT

```

```

v.product_id,
COUNT(DISTINCT v.user_id) AS viewers,
COUNT(DISTINCT a.user_id) AS adders,
COUNT(DISTINCT p.user_id) AS purchasers,
ROUND(100.0 * COUNT(DISTINCT p.user_id) / NULLIF(COUNT(DISTINCT v.user_id),0),
2) AS conv_rate_from_view
FROM views v
LEFT JOIN adds a ON a.user_id = v.user_id AND a.product_id = v.product_id AND
a.event_time BETWEEN v.event_time AND v.event_time + INTERVAL '7 days'
LEFT JOIN purchases p ON p.user_id = v.user_id AND p.product_id = v.product_id AND
p.event_time BETWEEN v.event_time AND v.event_time + INTERVAL '7 days'
GROUP BY v.product_id
ORDER BY conv_rate_from_view DESC NULLS LAST
LIMIT 100;

```

```

-----
-- 15) Security-conscious examples: RBAC hints (showcase understanding of permissions)
-----

```

```

-- create a read-only role for analytics consumers

DO $$

BEGIN

    CREATE ROLE analytics_reader NOLOGIN;

EXCEPTION WHEN duplicate_object THEN

    -- ignore

END $$;

```



```
GRANT USAGE ON SCHEMA prod_analytics TO analytics_reader;
```

```
GRANT SELECT ON ALL TABLES IN SCHEMA prod_analytics TO analytics_reader;
```

```
ALTER DEFAULT PRIVILEGES IN SCHEMA prod_analytics GRANT SELECT ON  
TABLES TO analytics_reader;
```

```
-----  
-- 16) Example show-casing schema evolution & migrations (safe ALTER patterns)
```

```
-----  
-- Add a new column with minimal locking (add column with NULL and then backfill in  
batches)
```

```
ALTER TABLE prod_analytics.events ADD COLUMN IF NOT EXISTS device_type TEXT;
```

```
-- Backfill in small batches (example pattern) to avoid long locks:
```

```
-- UPDATE prod_analytics.events SET device_type = (event_props->>'device') WHERE  
device_type IS NULL LIMIT 10000;
```

```
-----  
-- 17) Demonstration of advanced joins: anti-join, semi-join, lateral correlated subquery
```

```
-----  
-- Find users who never purchased (anti-join)
```

```
SELECT u.user_id, u.email, u.signup_ts
```

```
FROM prod_analytics.users u
```

```
LEFT JOIN prod_analytics.events e ON e.user_id = u.user_id AND e.event_type = 'purchase'
```

```
WHERE e.event_id IS NULL
```

```
LIMIT 100;
```

```
-- Semi-join: users who have at least one 'product_view'
```

```
SELECT u.user_id, u.email
FROM prod_analytics.users u
WHERE EXISTS (
    SELECT 1 FROM prod_analytics.events e WHERE e.user_id = u.user_id AND e.event_type =
'product_view'
)
LIMIT 100;
```

---

```
-- 18) Example: approximate distinct counts for cardinality estimation (hyperloglog via
extension)
```

---

```
CREATE EXTENSION IF NOT EXISTS postgresql_hll; -- if installed in environment
-- pseudo usage (may not be available in all envs): hll_add_agg & hll_cardinality
-- SELECT org_id, hll_cardinality(hll_add_agg(hll_hash_integer(user_id))) FROM
prod_analytics.events GROUP BY org_id;
```

---

```
-- 19) Sample test data insertion (kept minimal here; in a real test, swap with generated dataset)
```

---

```
INSERT INTO prod_analytics.organizations (org_name, industry)
VALUES ('Acme Corp', 'ecommerce'), ('Globex', 'cloud');
```

```
INSERT INTO prod_analytics.users (org_id, email, username, attrs)
SELECT o.org_id, format('user%s@example.com', gs), format('user%s', gs),
jsonb_build_object('signup_source', 'email')
FROM generate_series(1,100) gs
```

```
CROSS JOIN LATERAL (SELECT org_id FROM prod_analytics.organizations ORDER BY
random() LIMIT 1) o;
```

```
-- insert synthetic events (small sample)
```

```
INSERT INTO prod_analytics.events (user_id, org_id, event_type, event_props, event_time)
```

```
SELECT u.user_id, u.org_id,
```

```
    CASE WHEN random() < 0.02 THEN 'purchase' WHEN random() < 0.2 THEN
'add_to_cart' ELSE 'product_view' END,
```

```
    jsonb_build_object('price', round(random()*100,2), 'product_id', (1000 +
(random()*100)::int)::text, 'device', CASE WHEN random() < 0.5 THEN 'mobile' ELSE
'desktop' END),
```

```
    now() - (random() * interval '30 days')
```

```
FROM prod_analytics.users u
```

```
ORDER BY random()
```

```
LIMIT 5000;
```

```
-----
-- 20) Example README-style queries to demonstrate you understand trade-offs
-----
```

```
-- EXPLAIN (ANALYZE, BUFFERS, VERBOSE) <complex_query_here>;
```

```
-- For very large tables: prefer BRIN indexes, partition-pruning, and clustered order for HOT
updates.
```

```
-- Use VACUUM and ANALYZE regularly; use autovacuum tuning for high-throughput
ingestion.
```

```
-- For JSONB heavy workloads, GIN jsonb_path_ops is superior for key-existence queries; use
expression indexes for frequent lookups.
```

```
-----
```

-- 21) Cleanup helper (safe drops for exercise / demos)

-----

-- DROP SCHEMA prod\_analytics CASCADE; -- uncomment to teardown demo

-----

-- End of script: highlights & keywords for HR copy/paste

-----

-- HIGHLIGHTS:

-- \* Partitioning (RANGE by date) for high throughput time-series data

-- \* Indexing strategies: BTREE, BRIN, GIN (JSONB & full-text), trigram for fuzzy search

-- \* Advanced SQL constructs: window functions, recursive CTEs, lateral joins

-- \* Idempotent ETL patterns (staging + upsert ON CONFLICT)

-- \* Materialized views + concurrent refresh for expensive aggregates

-- \* PL/pgSQL ETL job examples with advisory locks for single-run concurrency

-- \* Observability: pg\_stat\_statements example and explain analyze pattern

-- Keywords: data engineering, partitioning, query optimization, window functions, JSONB, full-text search, upsert, CDC, ETL, materialized view, indexing, BRIN, GIN, PL/pgSQL, concurrency.